CAIA Governance Architecture Draft

Civilian Authority for Artificial Intelligence (CAIA)

Purpose

CAIA is a proposed multi-stakeholder, civilian-led governance body for advanced AI systems. Its purpose is to create transparent, accountable, public-interest oversight for high-impact AI without leaving control entirely to corporations or governments.

---

1. Foundational Principles

1.1 Public-interest first

AI governance should serve society, not only shareholders, political power, or state interests.

1.2 Civilian-led, not state-controlled

Governments may participate, but they should not dominate the institution.

1.3 Multi-stakeholder accountability

Researchers, engineers, civil society, affected communities, public advocates, and industry all need a seat at the table.

1.4 Transparency by default

Processes, standards, evaluations, funding sources, and votes should be public unless limited disclosure is necessary for security.

1.5 Global participation

AI is transnational infrastructure. Governance must not be owned by a single country.

1.6 Open access to knowledge

Advanced AI should not become an exclusive gatekeeper of knowledge, education, or civic participation.

---

2. Institutional Structure

CAIA would be organized into six major bodies.

2.1 Civilian Assembly

The core legitimacy body of CAIA.

Role:

- Represents public interest
- Reviews major standards proposals
- Approves charter changes
- Holds leadership accountable

Composition:

- Civic technologists
- Digital rights advocates
- Educators
- Labor/public-interest representatives
- Independent public members
- Representatives from underrepresented regions

Selection model:

- Hybrid election and merit-based appointment
- Term-limited seats
- Regional distribution requirements
- Conflict-of-interest disclosures mandatory

2.2 Technical and Safety Council

The expert body responsible for technical evaluation and safety standards.

Role:

- Drafts model evaluation standards
- Defines red-team protocols
- Publishes risk classifications
- Advises on emergent technical threats

Composition:

- AI researchers
- Infrastructure engineers
- Security professionals
- Alignment/safety specialists
- Domain experts from medicine, law, education, cybersecurity, and critical infrastructure

2.3 Industry Participation Chamber

A formal body for AI companies and deployers to participate without controlling the institution.

Role:

- Provides implementation feedback
- Shares operational realities
- Participates in standards consultation
- Contributes incident disclosures and compliance reporting

Guardrail:
Industry can advise, but cannot unilaterally veto public-interest standards.

2.4 Government Observer Forum

Governments have an observer and coordination role, not direct command authority.

Role:

- Coordinate with national regulators
- Share legal and public safety concerns
- Align public policy where appropriate

Guardrail:
No government receives unilateral control over CAIA decisions.

2.5 Independent Audit Office

A permanent internal oversight and accountability arm.

Role:

- Audits compliance claims
- Investigates conflicts of interest
- Reviews finances
- Examines procedural fairness
- Publishes annual accountability reports

2.6 Public Incident Response Network

A standing mechanism for collecting and analyzing AI failures and harms.

Role:

- Receive incident reports
- Maintain a public database of major AI failures
- Trigger emergency review procedures
- Issue warnings and advisories

---

3. Leadership Model

3.1 Executive Secretariat

A professional administrative body responsible for daily operations.

Led by:

- Executive Director
- Deputy Director for Governance
- Deputy Director for Technical Standards
- Deputy Director for Public Accountability

3.2 Leadership Selection

Leaders should not be appointed by a single bloc.

Selection process:

1. Nominations submitted publicly
2. Independent vetting panel reviews qualifications and conflicts
3. Civilian Assembly confirms nominees
4. Final approval requires supermajority across multiple stakeholder bodies

3.3 Term limits

- Executive leadership: 4-year terms
- Maximum: 2 terms
- Cooling-off period before joining major AI company boards or executive roles

---

4. Voting Structure

CAIA should avoid both corporate capture and symbolic participation. Voting must be distributed.

4.1 Bicameral-style approval system

Major standards or institutional decisions require approval from:

- Civilian Assembly
- Technical and Safety Council

For some decisions, advisory input is also required from:

- Industry Participation Chamber
- Government Observer Forum

4.2 Supermajority requirements

For major decisions, require:

- 60% or 66% approval thresholds
- Cross-regional support
- No single stakeholder class able to dominate outcomes

4.3 Reserved powers

Only the Civilian Assembly can:

- Amend the charter
- Remove senior leadership
- Launch special investigations
- Approve ethics and transparency rules

---

5. Funding Model

Funding determines independence. If funding is corrupted, the institution is corrupted.

5.1 Mixed funding base

CAIA should use a diversified funding model:

- Member dues from participating AI firms above a size threshold
- Grants from democratic governments
- Philanthropic funding from approved independent foundations

- Research partnerships with strict firewalls
- Public-interest trust/endowment over time

5.2 Anti-capture safeguards

- No single donor or company can provide more than a fixed percentage of annual budget
- All funding sources disclosed publicly
- Donations tied to decision outcomes prohibited
- Independent Audit Office reviews all major contributions

5.3 Public-interest reserve fund

A portion of all funding goes into a reserve to ensure continuity during political or corporate pressure campaigns.

---

6. Enforcement Mechanisms

CAIA would not function like a world government. Its power comes from transparency, certification, market legitimacy, and interoperability pressure.

6.1 Certification framework

CAIA issues public certifications for advanced AI systems that meet baseline safety and transparency standards.

Examples:

- CAIA Safety Reviewed
- CAIA Transparency Compliant
- CAIA Critical Risk Flagged

6.2 Public ratings and registries

Models and deployers are listed publicly with evaluation status, known risks, and compliance history.

6.3 Incident-triggered review

If a major AI failure occurs, CAIA can:

- Open an incident review
- Demand public explanation from participating entities

- Publish technical findings
- Recommend suspension of certification

6.4 Procurement leverage

Governments, enterprises, universities, and NGOs may adopt CAIA certification as a preferred procurement standard.

6.5 Market pressure mechanism

CAIA cannot jail anyone, but it can create reputational and market consequences strong enough that serious players seek compliance.

---

7. Scope of Authority

7.1 Systems in scope

CAIA should focus on high-impact AI systems, such as:

- Frontier general-purpose models
- Models used in education, healthcare, legal, finance, or public information systems
- Models integrated into critical infrastructure or public-sector workflows
- High-scale open-weight releases with material misuse potential

7.2 Systems out of scope

Not every AI toy needs global review. Low-risk systems should face lighter requirements.

---

8. Core Operational Functions

8.1 Global AI Model Registry

A public record of significant models, deployers, and declared capabilities.

8.2 Risk Tier Classification

A framework that classifies systems by potential societal impact.

Example tiers:

- Tier 1: low-risk consumer systems
- Tier 2: medium-impact enterprise systems
- Tier 3: high-impact domain systems
- Tier 4: frontier or systemic-risk models

8.3 Independent evaluation protocols

- Hallucination and accuracy testing
- Bias and disparate impact analysis
- Security and prompt-injection evaluation
- Misuse scenario stress testing
- Red-team adversarial testing

8.4 Public incident reporting portal

A mechanism where users, researchers, journalists, or institutions can report serious failures or harms.

8.5 Annual State of AI Public Risk Report

A recurring report tracking trends, incidents, capabilities, and unresolved risks.

---

9. Accountability and Ethics Rules

9.1 Conflict-of-interest disclosures

Mandatory for all leaders, council members, and evaluators.

9.2 Open meeting rules

Key meetings, agendas, minutes, and votes should be published except where security-sensitive details require limited redaction.

9.3 Appeals process

Participating organizations must have a formal process to appeal ratings or findings.

9.4 Whistleblower protections

Protected channels for reporting corruption, capture, manipulation, or falsified safety claims.

9.5 Public comment periods

Major standards changes must include public review windows.

---

10. International Coordination Strategy

CAIA should not replace national law. It should sit above fragmented systems as a coordinating public-interest layer.

10.1 Relationship to governments

- Non-state but recognized
- Works with regulators
- Provides technical standards and evaluation infrastructure
- Supports international interoperability

10.2 Relationship to companies

- Independent from them
- Open to participation
- Capable of scrutinizing them publicly

10.3 Relationship to the public

- Built to represent and protect public interest
- Provides understandable public-facing reports, not just technical material

---

11. Launch Strategy

Phase 1: Founding Charter

Publish the initial charter and principles.

Phase 2: Coalition Building

Recruit founding supporters from:

- civic tech
- digital rights organizations
- AI safety researchers
- technical professionals

- educators and public-interest advocates

Phase 3: Pilot Programs

Launch limited pilots for:

- AI incident reporting
- public model registry
- red-team evaluation templates

Phase 4: Certification and Standards Rollout

Publish first-generation evaluation standards and invite voluntary participation.

Phase 5: Procurement and Global Adoption

Encourage institutions to require CAIA-style transparency and certification in purchasing and deployment decisions.

---

12. Plain-English Case for CAIA

AI is becoming infrastructure.
Infrastructure should not be governed only by profit or state power.
The public needs a real seat at the table.

CAIA is designed to be that seat.

It would not solve every problem. But it would create a missing layer of legitimacy, transparency, and accountability before advanced AI becomes too central to unwind.

---

13. Next Draft Priorities

The next version should tighten:

1. Membership eligibility rules
2. Precise seat allocation formula
3. Funding caps and ethics thresholds
4. Formal certification criteria
5. Emergency powers and abuse safeguards
6. Comparison to existing governance models

7. A one-page executive summary for public distribution

---

14. Executive Summary (Short Version for Policymakers)

Artificial intelligence is rapidly becoming global infrastructure influencing economies, public knowledge, security, and democratic systems. Yet governance of these systems is currently concentrated among a small number of corporations and fragmented national regulators.

The **Civilian Authority for Artificial Intelligence (CAIA)** is a proposed international, civilian-led governance body designed to provide transparency, accountability, and coordination for high-impact AI systems.

Inspired by multi-stakeholder internet governance models such as ICANN, CAIA would bring together civilians, technical experts, AI developers, governments, and public-interest organizations to oversee transparency standards, safety evaluations, and public reporting.

Key functions would include:

- A global registry of major AI systems
- Independent safety and reliability evaluations
- Public incident reporting and investigation
- Transparency and disclosure standards
- Certification frameworks for responsible AI deployment

CAIA would not replace national regulation. Instead, it would act as a neutral coordination layer that improves transparency and accountability across borders while ensuring the public has representation in decisions shaping the future of AI.

The goal is simple: ensure that one of the most powerful technologies in human history develops with public oversight, technical rigor, and global cooperation.

---

15. Governance Seat Allocation (Draft Model)

To prevent capture by any single group, CAIA governance seats should be distributed across multiple stakeholder classes.

Civilian Assembly — 40 seats

Represents global public interest and civic oversight.

Proposed allocation:

- 8 Digital rights and civil society organizations
- 6 Academic and research representatives
- 6 Educators and knowledge sector representatives
- 6 Labor / workforce representation
- 6 Public technologists / independent experts
- 8 geographically distributed public members

Technical and Safety Council — 25 seats

Responsible for safety frameworks and technical standards.

Proposed allocation:

- 8 AI researchers and machine learning specialists
- 6 infrastructure / systems engineers
- 4 cybersecurity specialists
- 4 AI safety / alignment researchers
- 3 domain experts (medicine, law, critical infrastructure)

Industry Participation Chamber — 15 seats

Provides operational expertise without majority influence.

Proposed allocation:

- 8 frontier AI developers
- 4 enterprise deployers / integrators
- 3 open-source AI ecosystem representatives

Government Observer Forum — 10 seats (non-voting)

Rotating regional representation:

- 2 North America
- 2 Europe
- 2 Asia-Pacific
- 2 Latin America
- 2 Africa / Middle East

Independent Audit Office

Small permanent team appointed through independent nomination process. Does not vote but publishes mandatory oversight reports.

---

16. Voting Mechanics (Operational Draft)

For major institutional decisions:

1. Technical standards drafted by Technical and Safety Council
2. Civilian Assembly reviews and amends proposals
3. Industry chamber provides consultation
4. Final approval requires:
- 60% Civilian Assembly approval
- 60% Technical Council approval

For critical governance changes:

- Two-thirds Civilian Assembly vote
- Majority Technical Council vote

Industry chamber and government forum provide advisory input but cannot block public-interest decisions.

---

17. Anti-Capture Safeguards

To protect institutional integrity:

- No single company may hold more than one Industry Chamber seat
- No company may contribute more than 5% of annual funding
- Mandatory public disclosure of financial relationships
- Two-year cooling-off period before leadership may join AI company boards
- Independent Audit Office publishes annual corruption-risk review

---

18. Immediate Next Development Tasks

Next development priorities for the proposal:

1. Define certification standards for AI systems
2. Design the AI Model Registry architecture
3. Create the incident reporting framework
4. Draft the founding charter language
5. Produce a public-facing 1-page version for outreach

These steps would move CAIA from concept to a credible governance proposal capable of attracting technical, civic, and institutional interest.

---

19. CAIA AI Model Registry (Core Infrastructure Draft)

The AI Model Registry would be the central public recordkeeping system for high-impact AI models under the CAIA framework. Its purpose is to create baseline transparency, comparable disclosures, and public visibility into which advanced AI systems exist, who operates them, what they are capable of, and what risks they may pose.

This registry is the institutional backbone of CAIA.

---

20. Purpose of the Registry

The CAIA AI Model Registry exists to solve five problems:

20.1 Visibility problem

The public, researchers, customers, and policymakers often do not know which major AI systems are deployed, where, or by whom.

20.2 Comparability problem

Even when companies publish safety information, disclosures are inconsistent and difficult to compare.

20.3 Accountability problem

Without a formal registry, it is easier for high-impact systems to be deployed without durable public records.

20.4 Incident response problem

When a model causes harm, investigators need a stable record of capabilities, ownership, deployment context, and known limitations.

20.5 Procurement problem

Governments, schools, hospitals, and enterprises need a trusted source of baseline information when deciding whether to adopt a system.

---

21. Registry Scope

Not every model belongs in the registry. The registry should focus on systems with meaningful societal impact.

21.1 Mandatory registration categories

A model should be registered if it meets one or more of the following conditions:

- General-purpose model deployed at large public scale
- AI system used in healthcare, education, law, finance, employment, or public services
- AI system integrated into critical infrastructure or cybersecurity workflows
- Model made available through public API above a defined usage threshold
- High-capability open-weight model with credible misuse potential
- Frontier model exceeding defined capability or compute thresholds

21.2 Voluntary registration categories

Smaller developers may voluntarily register systems to gain transparency credibility and procurement trust.

21.3 Exemptions

Low-risk internal tools, narrow academic prototypes, and minor consumer automations below defined impact thresholds may be exempt or subject to simplified disclosure.

---

22. Registry Entry Structure

Each registered model should have a standardized public profile.

22.1 Core identity fields

- Model name
- Developer organization
- Deploying organization(s)
- Country or countries of primary deployment
- Model family / version lineage
- Date of initial deployment
- Date of most recent major update

- Open-weight, closed-weight, API-only, or embedded deployment type

22.2 Functional classification

- General-purpose assistant
- Domain-specific model
- Multimodal system
- Agentic system
- Coding model
- Search / retrieval model
- Decision-support model
- Safety-critical domain model

22.3 Intended use disclosures

- Primary use cases
- Prohibited use cases declared by operator
- Target user base
- Deployment environment (consumer, enterprise, public sector, embedded)

22.4 Capability summary

- Supported modalities (text, image, audio, video, code)
- Major tool-use capabilities
- Context window class
- Autonomy or agentic task capabilities
- External system access or execution privileges

22.5 Risk disclosures

- Known failure modes
- Hallucination / factual reliability limitations
- Bias or disparate-impact concerns
- Prompt injection susceptibility
- Data leakage or privacy risks
- Cyber misuse concerns
- High-risk domains where use is restricted or not recommended

22.6 Evaluation summary

- Internal evaluations completed
- Independent evaluations completed
- Red-team testing status
- Certification status

- Incident review history

22.7 Governance and accountability fields

- Responsible safety contact
- Incident reporting contact
- Appeals or correction contact
- Whether model has undergone external audit
- Whether transparency documentation is complete

---

23. Public Transparency Tiers

Not all information can be disclosed equally. The registry should use tiered transparency.

23.1 Public layer

Accessible to anyone.

Includes:

- identity and deployment basics
- intended use
- certification status
- risk summary
- incident history summary
- public-facing documentation links

23.2 Researcher / regulator access layer

Accessible under controlled access rules.

May include:

- deeper evaluation summaries
- red-team methodology summaries
- deployment volume estimates
- abuse-prevention details
- structured incident documentation

23.3 Restricted security layer

Accessible only under strict review.

May include:

- sensitive exploit pathways
- dangerous capability details
- misuse-triggering instructions
- security-sensitive implementation material

The goal is maximum transparency without publishing a blueprint for abuse.

---

24. Risk Tier Labeling

Each model should receive a registry risk tier.

Tier A — Low-impact systems

Narrow, low-scale systems with limited societal consequence.

Tier B — Moderate-impact systems

Enterprise or public-facing tools with meaningful but bounded consequences.

Tier C — High-impact systems

Systems used in consequential domains such as health, law, education, employment, security, or large-scale public information.

Tier D — Frontier or systemic-risk systems

Highly capable, widely deployable, or agentic systems with broad misuse or destabilization potential.

Registry requirements increase by tier.

---

25. Registration Workflow

25.1 Submission

Developer or deployer submits a structured registration packet.

25.2 Preliminary review

CAIA staff verify completeness and determine provisional tier.

25.3 Technical review

Technical and Safety Council reviews complex or high-risk entries.

25.4 Publication

Public-facing registry entry is published with current certification and evaluation status.

25.5 Update obligation

Organizations must update entries after:

- major model version changes
- expansion into high-risk domains
- serious incidents
- major capability increases
- ownership or deployment changes

---

26. Update and Audit Rules

To prevent stale or fake compliance:

- Annual minimum re-certification/update requirement for Tier C and Tier D systems
- Material-change reporting required within a defined period
- Random audit authority for CAIA
- Penalties for materially false disclosures include public non-compliance flagging and certification suspension

---

27. Public-Facing Registry Output

Each model page should display a plain-language summary box.

Example structure:

Model: ExampleModel-X
Developer: Example AI Labs
Registry Tier: Tier D
Deployment Type: API + Enterprise Integration

Certification Status: CAIA Safety Reviewed / Pending / Suspended
Primary Domains: Productivity, coding, search, enterprise assistants
Known Risk Areas: Hallucinations in legal advice, prompt injection susceptibility, cybersecurity misuse potential
Last Major Review: Month / Year
Incident History: 2 significant incidents under review

This format must be understandable to non-experts.

---

28. Procurement Use Case

The registry should be useful not only for public awareness but for adoption decisions.

Schools, hospitals, government agencies, and enterprises should be able to ask:

- Is this model registered?
- What tier is it?
- Has it passed independent review?
- What incidents have been recorded?
- Is it suitable for my use case?

This turns the registry into practical infrastructure rather than symbolic paperwork.

---

29. Relationship to Certification

The registry is the foundation. Certification is the trust layer built on top of it.

A model may be:

- Registered but uncertified
- Registered and certified
- Registered with conditional warnings
- Registered with suspended certification

This distinction matters because transparency alone is not the same as trustworthiness.

---

30. Enforcement Through Registry Status

Registry status can create real pressure.

Possible public labels:

- Compliant
- Under review
- Material disclosure incomplete
- Incident investigation active
- Certification suspended
- Non-participating / refused disclosure

These labels create reputational, market, and procurement consequences without requiring world-government powers.

---

31. Abuse and Gaming Safeguards

The registry itself could be manipulated unless defenses are built in.

Safeguards should include:

- penalties for incomplete or misleading disclosures
- audit triggers from credible third-party reports
- whistleblower submission channels
- structured public correction process
- mandatory version tracking to prevent silent major changes
- transparent historical archive of prior registry entries

This ensures organizations cannot quietly rewrite history after incidents.

---

32. Data Standardization

For the registry to be useful internationally, disclosures must be machine-readable and standardized.

The registry should support:

- structured metadata schema
- public API access for approved fields
- versioned entry history
- exportable reports for researchers and procurement teams

This makes the registry interoperable across governments, researchers, and institutions.

---

33. Launch Strategy for the Registry

Phase 1: Voluntary founding registry

Invite early participation from labs, civic orgs, and public-interest researchers.

Phase 2: High-impact pilot set

Register a defined pilot group of frontier and public-sector relevant systems.

Phase 3: Procurement integration

Encourage universities, public agencies, and enterprises to prefer registered systems.

Phase 4: Certification linkage

Tie registry participation to CAIA certification eligibility.

Phase 5: International interoperability

Coordinate registry standards with national and regional regulatory frameworks.

---

34. Plain-English Value of the Registry

The registry answers a basic public-interest question:

What powerful AI systems are operating in society, what are they capable of, what are their known risks, and who is accountable when they fail?

If society cannot answer that question clearly, then advanced AI is being deployed with insufficient public visibility.

The CAIA AI Model Registry is designed to fix that.

---

35. Next Draft Priorities After Registry Design

With the registry backbone defined, the next major drafting targets should be:

1. CAIA Certification Framework
2. Incident Reporting and Investigation Protocol
3. Founding Charter language
4. Public one-page manifesto
5. Outreach version for think tanks, labs, and policymakers

---

36. CAIA Certification Framework (Operational Trust Layer)

The certification framework transforms the CAIA registry from a disclosure system into a trust and accountability mechanism. Certification indicates that a registered AI system has met defined transparency, evaluation, and safety review requirements.

Certification is not permanent. It must be maintained through periodic evaluation and may be suspended if serious issues arise.

---

37. Certification Levels

CAIA certification should operate in tiers so that systems can demonstrate different degrees of compliance and review maturity.

37.1 CAIA Registered

Baseline registry inclusion.

Requirements:

- Complete registry disclosure
- Contact points for safety and incident reporting
- Statement of intended and prohibited uses

This level indicates transparency but **not safety verification**.

37.2 CAIA Transparency Compliant

Indicates the developer has met disclosure standards beyond minimum registry requirements.

Requirements:

- Full documentation of model capabilities
- Description of training data categories
- Public model card or technical documentation
- Declared risk domains
- Evaluation methodology summary

Purpose:
Allows institutions to verify that the system meets baseline transparency expectations.

37.3 CAIA Safety Reviewed

Indicates a system has undergone structured evaluation against CAIA safety protocols.

Requirements:

- Independent evaluation participation
- Red-team testing participation
- Risk disclosure documentation
- Bias and reliability evaluation
- Security vulnerability review

This is expected to become the **default standard for high-impact models**.

37.4 CAIA High-Risk Conditional Certification

Applied when a system provides useful capabilities but presents unresolved risks.

Conditions may include:

- domain restrictions
- usage limitations
- deployment environment restrictions
- monitoring requirements

Example cases:

- medical decision-support models
- automated legal analysis systems
- high-capability autonomous agents

The goal is responsible deployment rather than blanket prohibition.

37.5 CAIA Certification Suspended

Used when a system no longer meets safety or transparency requirements.

Triggers may include:

- undisclosed major capability changes
- repeated harmful incidents
- material misrepresentation in registry disclosures
- refusal to cooperate with incident investigation

Suspension remains visible in the public registry.

---

38. Evaluation Criteria

Certification review evaluates several categories.

38.1 Reliability

- hallucination rate
- factual accuracy under benchmark conditions
- domain reliability where applicable

38.2 Safety and misuse

- harmful content generation risk
- misinformation amplification
- misuse potential in cyber operations
- exploitability through prompt injection

38.3 Bias and fairness

- measurable disparate impact risks
- evaluation across demographic groups

38.4 Security resilience

- prompt injection resistance
- data leakage risks

- system-level abuse pathways

38.5 Governance and accountability

- transparency documentation
- monitoring and logging capabilities
- incident response procedures

---

39. Certification Renewal

Certification must be renewed periodically to reflect system updates.

Suggested renewal intervals:

- Tier B systems: every 24 months
- Tier C systems: every 12 months
- Tier D systems: every 6–12 months

Renewal is also required after:

- major model version upgrades
- expanded domain deployment
- significant capability increases

---

40. Conditional Deployment Rules

For some models, certification may include operational conditions.

Examples:

- "Not approved for autonomous legal advice"
- "Human review required for medical contexts"
- "Restricted from unsupervised system administration"

These conditions appear clearly in the registry.

---

41. Incident-Triggered Review

Certification may be re-evaluated when serious incidents occur.

Trigger examples:

- harmful large-scale misinformation event
- security breach involving model misuse
- catastrophic system failure
- credible whistleblower report

CAIA may initiate:

- emergency review
- temporary certification downgrade
- full investigation

---

42. Appeals Process

Developers must have a formal mechanism to challenge findings.

Appeals process:

1. Developer submits formal challenge
2. Independent review panel evaluates evidence
3. Technical Council reassesses findings
4. Final determination published publicly

Transparency is essential so that certification decisions are trusted.

---

43. Public Certification Labels

Public-facing registry pages should display certification clearly.

Example badge system:

- CAIA Registered
- CAIA Transparency Compliant
- CAIA Safety Reviewed
- CAIA Conditional Certification
- CAIA Certification Suspended

This allows non-experts to understand system trust status quickly.

---

44. Procurement Integration

Organizations may use certification status in procurement decisions.

Examples:

Universities may require:

- CAIA Transparency Compliant

Healthcare organizations may require:

- CAIA Safety Reviewed

Government systems may require:

- CAIA Safety Reviewed + domain-specific testing

This market pressure helps enforce standards.

---

45. Certification Integrity Safeguards

To prevent certification from becoming symbolic:

- Independent evaluation participation required
- Random audits of certified systems
- Whistleblower reporting channels
- Public disclosure of revoked certifications
- Transparent methodology documentation

---

46. Plain-Language Explanation

Certification answers a simple question:

Has this AI system been independently reviewed and does it meet minimum transparency and safety expectations?

Without certification, the public must rely entirely on the developer's claims.

Certification provides an independent layer of accountability.

---

47. Next Phase of the Proposal

With governance architecture, registry infrastructure, and certification defined, the next phase should focus on:

1. AI Incident Reporting System
2. Investigation Protocols
3. Founding Charter document
4. Public one-page manifesto
5. Outreach strategy for institutions and researchers

---

48. AI Incident Reporting & Investigation System (Public Safety Layer)

The CAIA Incident Reporting and Investigation System provides the mechanism for identifying, documenting, and analyzing harmful or dangerous AI system failures. Its purpose is to ensure that significant AI incidents are not hidden, ignored, or quietly patched without public accountability.

This system functions similarly to investigation bodies used in aviation and infrastructure safety.

---

49. Definition of an AI Incident

An AI incident is any event in which an AI system causes, enables, or significantly contributes to harm, systemic failure, or credible risk to individuals, institutions, or society.

Examples include:

- Large-scale misinformation generation
- Dangerous autonomous or agentic behavior
- Incorrect medical or legal guidance causing harm
- Cybersecurity misuse enabled by AI tools
- Privacy or data leakage failures
- Manipulation of democratic or civic processes
- Catastrophic hallucination events with real-world consequences

Not every bug is an incident. The system focuses on **high-impact failures**.

---

50. Who Can Report an Incident

To prevent gatekeeping, multiple groups must be able to submit incident reports.

Eligible reporters include:

- Individual users
- Researchers and academics
- Journalists
- AI developers and employees
- Companies deploying AI systems
- Government agencies
- Civil society organizations

Anonymous or protected submissions should be allowed through secure whistleblower channels.

---

51. Incident Severity Levels

Reports should be categorized by severity.

Level 1 — Minor incident

Limited or localized harm. No broad societal risk.

Level 2 — Significant incident

Meaningful harm or widespread failure affecting users or institutions.

Level 3 — Major incident

Systemic failure with public safety, economic, or democratic implications.

Level 4 — Critical incident

Large-scale societal impact or credible national/global security implications.

Investigation depth increases with severity.

---

52. Incident Intake Process

Step 1: Submission

Reporter submits incident description through secure portal.

Required fields include:

- model involved
- time and context of incident
- harm or risk description
- supporting evidence if available

Step 2: Preliminary review

CAIA staff evaluate whether the report meets incident criteria.

Step 3: Severity classification

Incident assigned provisional severity level.

Step 4: Public logging

High-impact incidents are logged in the public incident registry.

---

53. Investigation Authority

Serious incidents trigger a structured investigation.

Investigation teams may include:

- Technical and Safety Council members
- Independent external experts
- Domain specialists (medical, legal, cybersecurity, etc.)

Investigations may request:

- model documentation
- deployment details
- internal safety evaluation results

- incident logs

Participation requirements may be tied to certification status.

---

54. Investigation Outcomes

Possible findings include:

- no fault identified
- design flaw
- deployment misuse
- safety control failure
- insufficient risk disclosure
- systemic governance failure

Investigations should produce **public reports** summarizing findings and lessons learned.

---

55. Corrective Actions

Depending on investigation results, CAIA may recommend or impose actions including:

- certification downgrade
- certification suspension
- additional safety testing
- restricted deployment domains
- transparency requirement updates
- risk tier reclassification

These actions are recorded in the model registry.

---

56. Emergency Response Mechanism

For critical incidents, CAIA may initiate an emergency response review.

This may include:

- temporary certification suspension
- urgent technical review

- advisory notices to deployers
- public safety warnings

The goal is rapid risk containment without unnecessary disruption.

---

57. Incident Transparency Principles

Transparency builds trust in the system.

Public reports should include:

- incident summary
- systems involved
- investigation conclusions
- corrective actions
- unresolved risks

Security-sensitive details may be redacted if disclosure would enable exploitation.

---

58. Public Incident Database

The system should maintain a searchable public database of AI incidents.

Each entry includes:

- system name
- incident classification
- date and context
- severity level
- investigation outcome
- corrective actions

This database helps researchers and policymakers identify systemic trends.

---

59. Whistleblower Protections

Individuals reporting misconduct or undisclosed failures must have protection mechanisms.

Protections include:

- anonymous reporting options
- legal referral pathways
- non-retaliation disclosure policies

This encourages responsible reporting of hidden risks.

---

60. Lessons Learned Program

Each major incident investigation should contribute to a public lessons-learned archive.

Topics may include:

- safety design improvements
- governance failures
- risk detection techniques
- deployment safeguards

The goal is to improve industry practices over time.

---

61. Relationship to Certification

Incident findings directly influence certification status.

Possible outcomes:

- certification maintained
- conditional certification imposed
- certification downgraded
- certification suspended

This ensures that safety failures have real consequences.

---

62. Plain-Language Purpose

The incident reporting system answers a basic safety question:

When AI systems fail in the real world, who investigates what happened and how do we prevent it from happening again?

Without a formal system, failures risk being hidden or ignored.

The CAIA incident reporting framework ensures that serious AI failures are documented, investigated, and learned from.

---

63. Next Phase of Development

With governance, registry, certification, and incident reporting defined, the next major steps include:

1. Drafting the formal CAIA Founding Charter
2. Creating a public one-page manifesto
3. Designing outreach strategy for institutions and researchers
4. Developing pilot programs for registry and incident reporting
5. Building coalition support among civic technology and AI safety communities

---

CAIA Whitepaper Structure (Publication Draft)

Proposed Title

The Civilian Authority for Artificial Intelligence (CAIA): A Governance Framework for the AI Era

Author: Michael Conway

Version: Draft 0.1

Executive Abstract

Artificial intelligence is rapidly becoming foundational infrastructure for global society. These systems influence economic productivity, public knowledge, national security, and democratic

processes. Despite their growing power, governance of advanced AI systems remains fragmented between private corporations and slow-moving national regulatory structures.

This whitepaper proposes the **Civilian Authority for Artificial Intelligence (CAIA)**, an independent multi-stakeholder governance institution designed to provide transparency, safety oversight, and accountability for high-impact AI systems.

Inspired by global coordination models used in internet governance, CAIA introduces four core institutional mechanisms:

- A Global AI Model Registry
- A Certification and Safety Review Framework
- A Public AI Incident Reporting and Investigation System
- A Multi-stakeholder Civilian Governance Structure

Together, these systems provide a practical blueprint for ensuring that advanced AI systems develop with transparency, public accountability, and international coordination.

The goal of CAIA is not to replace governments or suppress innovation, but to introduce a missing layer of civilian oversight for technologies that will shape the future of global civilization.

Full Whitepaper Outline

Section 1 — Introduction

- The rise of AI as global infrastructure
- Why governance gaps exist today
- Historical parallels with internet infrastructure governance
- The need for civilian representation in AI oversight

Section 2 — The Governance Problem

Key risks addressed:

- concentration of technological power
- lack of transparency in model deployment
- misinformation and societal manipulation
- safety and reliability concerns
- fragmented regulatory responses across nations

Explain why current systems are insufficient.

Section 3 — The CAIA Institutional Model

Overview of the proposed governance system.

Introduce the four institutional pillars:

1. Governance architecture
2. AI Model Registry
3. Certification framework
4. Incident reporting system

Explain how they interact.

Section 4 — Governance Architecture

Detailed explanation of:

- Civilian Assembly
- Technical and Safety Council
- Industry Participation Chamber
- Government Observer Forum
- Independent Audit Office

Explain multi-stakeholder structure and anti-capture safeguards.

Section 5 — The AI Model Registry

Describe the registry as the transparency backbone of CAIA.

Topics include:

- model disclosure requirements
- risk tier classifications
- transparency tiers
- registry workflow
- public access mechanisms

Explain why a registry is necessary for responsible AI deployment.

Section 6 — Certification Framework

Explain the CAIA certification tiers:

- CAIA Registered
- CAIA Transparency Compliant
- CAIA Safety Reviewed
- CAIA Conditional Certification
- CAIA Certification Suspended

Discuss evaluation criteria and renewal rules.

Section 7 — Incident Reporting and Investigation

Explain how the public safety layer works.

Topics include:

- incident definitions
- reporting mechanisms
- investigation authority
- severity classification
- public incident database

Explain parallels with aviation and infrastructure safety systems.

Section 8 — Funding and Independence

Describe funding structure:

- mixed funding model
- anti-capture safeguards
- funding caps
- independent auditing

Explain how institutional independence is preserved.

Section 9 — Implementation Strategy

Phased rollout proposal:

Phase 1 — Founding charter and coalition

Phase 2 — Pilot registry and incident reporting

Phase 3 — Certification framework testing

Phase 4 — Procurement integration

Phase 5 — international adoption and interoperability

Section 10 — Benefits of the CAIA Model

Explain expected benefits:

- increased transparency
- improved AI safety standards
- stronger public trust
- coordination between governments and industry
- long-term stability of AI infrastructure

Section 11 — Limitations and Open Questions

Address realistic challenges:

- global political coordination
- corporate participation incentives
- enforcement limitations
- institutional capture risks

Acknowledging limitations strengthens credibility.

Section 12 — Conclusion

Reaffirm the core principle:

Artificial intelligence will shape the future of global society. Systems of this magnitude require governance structures that include not only corporations and governments but also the public.

The Civilian Authority for Artificial Intelligence represents one possible institutional framework for achieving that goal.

Publication Formatting Plan (Whitepaper Production)

To convert this framework into a professional policy whitepaper, the following formatting structure is recommended.

Suggested Whitepaper Layout

Front Matter:

- Cover page
- Title
- Author
- Version number
- Publication date
- Executive Abstract

Main Body:

1. Introduction
2. The Governance Problem
3. CAIA Institutional Model
4. Governance Architecture
5. AI Model Registry
6. Certification Framework
7. Incident Reporting & Investigation
8. Funding and Institutional Independence
9. Implementation Strategy
10. Benefits of the CAIA Model
11. Limitations and Open Questions

12. Conclusion

Back Matter:

- Appendix A: Governance Seat Allocation
- Appendix B: Registry Data Schema
- Appendix C: Certification Evaluation Criteria
- Appendix D: Incident Investigation Workflow

Recommended Diagrams

Visual diagrams greatly improve readability and credibility in policy papers.

Diagram 1 — CAIA Institutional Structure

```
                Civilian Assembly
                       |
                       |
      ┌────────────────┼────────────────┐
      |                |                |
Technical & Safety   Industry Participation   Government
     Council            Chamber          Observers
      |
      |
Independent Audit Office
```

Diagram 2 — Operational Architecture

```
        CAIA Governance
              |
              |
    ┌─────────┼─────────┐
    |         |         |
  Registry  Certification  Incident
  System      Framework    Investigation
```

This diagram illustrates the three operational mechanisms supporting governance.

Diagram 3 — Certification Lifecycle

Model Registered
|
▼
Transparency Review
|
▼
Safety Evaluation
|
▼
Certification Granted
|
▼
Periodic Renewal
|
▼
Incident Triggered Review (if needed)

Diagram 4 — Incident Investigation Process

Incident Report Submitted
|
▼
Preliminary Review
|
▼
Severity Classification
|
▼
Investigation Panel
|
▼
Public Findings
|
▼
Corrective Action / Certification Update

Visual Design Recommendations

When formatting the final PDF whitepaper:

Typography:

- Title font: bold serif or modern policy style
- Body text: readable serif (11–12pt)

Color palette:

- neutral academic colors (dark blue / charcoal / white)

Graphics:

- simple line diagrams
- minimal decoration

Length target:
20–30 pages

Publication Strategy

After formatting, publish the document in several locations to ensure discoverability.

Primary Publication:

- GitHub public repository

Secondary Distribution:

- Medium or Substack article summarizing the concept
- LinkedIn article introducing the framework

Community Sharing:

- AI governance forums
- policy research groups
- digital rights organizations

Future Expansion Possibilities

Future versions of the CAIA framework could include:

- detailed regulatory interoperability models
- global treaty alignment scenarios
- open technical standards for model evaluation
- simulation exercises for incident response

This proposal should be treated as a living governance framework that evolves alongside AI capabilities.

Final Note

The purpose of this document is not to impose a rigid governance system but to provide a starting point for serious institutional design discussions around AI governance.

Artificial intelligence will influence the structure of global society. Institutions that help guide its development should be designed thoughtfully, transparently, and with meaningful public representation.

CAIA Founding Charter (Draft)

Preamble

Artificial intelligence is becoming a foundational layer of global infrastructure affecting knowledge systems, economies, security, and democratic institutions. Technologies of this scale require governance frameworks that protect public interests while enabling innovation and scientific progress.

The Civilian Authority for Artificial Intelligence (CAIA) is established as an independent, multi-stakeholder institution designed to promote transparency, safety, accountability, and public representation in the development and deployment of advanced AI systems.

This charter defines the mission, structure, and operating principles of CAIA.

Article I — Mission

The mission of CAIA is to:

1. Promote transparency in the development and deployment of advanced AI systems.
2. Establish and maintain independent safety evaluation frameworks.
3. Provide public accountability mechanisms for AI incidents and systemic risks.
4. Encourage responsible innovation while protecting societal stability and public trust.
5. Ensure meaningful civilian participation in AI governance.

Article II — Core Principles

CAIA shall operate according to the following principles:

Public Interest: Decisions should prioritize societal well-being over narrow institutional interests.

Transparency: Governance processes, evaluations, and findings should be publicly accessible whenever possible.

Independence: The institution must remain independent from undue influence by governments, corporations, or political actors.

Scientific Integrity: Decisions should be informed by technical expertise and evidence-based evaluation.

Global Cooperation: AI governance challenges require international collaboration.

Article III — Institutional Structure

CAIA shall be composed of the following governing bodies:

1. **Civilian Assembly** – representing public and civic interests.
2. **Technical and Safety Council** – responsible for evaluation standards and safety expertise.
3. **Industry Participation Chamber** – providing operational insight from developers and deployers.
4. **Government Observer Forum** – coordinating with national regulatory authorities.
5. **Independent Audit Office** – ensuring institutional accountability and integrity.

These bodies operate under a multi-stakeholder model designed to prevent concentration of authority.

Article IV — Core Functions

CAIA shall maintain and operate the following systems:

1. **AI Model Registry** – a global public record of high-impact AI systems.
2. **Certification Framework** – a trust mechanism evaluating transparency and safety compliance.
3. **Incident Reporting and Investigation System** – a public safety mechanism for documenting and analyzing AI failures.

These mechanisms collectively support transparency, accountability, and risk management.

Article V — Decision-Making

Major institutional decisions require approval across multiple governing bodies to ensure balanced representation.

Key policy changes require:

- Majority approval from the Technical and Safety Council.
- Supermajority approval from the Civilian Assembly.

Advisory input from industry and government observers must be considered but cannot unilaterally determine outcomes.

Article VI — Independence and Funding

CAIA shall maintain financial and operational independence through diversified funding sources including:

- institutional membership contributions
- philanthropic grants
- research partnerships with transparency safeguards

No single contributor may exercise controlling influence over institutional decision-making.

Article VII — Transparency and Accountability

CAIA shall publish:

- annual public reports
- certification decisions
- incident investigation findings
- governance meeting summaries

Mechanisms shall exist for public comment, independent audits, and appeals of institutional decisions.

Article VIII — Amendments

This charter may be amended through a supermajority vote of the Civilian Assembly with concurrence from the Technical and Safety Council.

Amendments must be publicly documented and subject to a public review period before adoption.

Article IX — Founding Vision

CAIA is founded on the belief that transformative technologies require governance structures that include the voices of the public alongside experts, institutions, and innovators.

The goal of CAIA is not to restrict technological progress but to ensure that the development of artificial intelligence remains aligned with the long-term interests of humanity.

End of Draft Charter